

Investigating how well a unilateral dumbbell bicep curl was performed

Neil Caren

March 20, 2016

Summary

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement – a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it.

The goal of this reserach is to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways. More information is available from the website here: <http://groupware.les.inf.puc-rio.br/har> (<http://groupware.les.inf.puc-rio.br/har>) (see the section on the Weight Lifting Exercise Dataset).

Environment Setup and data acqision

The training data for this project are available here:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>
(<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>)

The test data are available here:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>
(<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>)

```
library(caret)
```

```
## Loading required package: lattice  
## Loading required package: ggplot2
```

```
library(randomForest)
```

```
## randomForest 4.6-12
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
##
## The following object is masked from 'package:ggplot2':
##
##     margin
```

```
setwd("~/Dropbox/scans/datascience/John Hopkins/Machine Learning")
training <- read.csv("pml-training.csv", stringsAsFactors = FALSE)
testing <- read.csv("pml-testing.csv", stringsAsFactors = FALSE)
```

Exploratory Data Analysis

The data has many variables with significant NA data points. After impact analysis it was decided to remove these columns along with other non-numerical variable that would get in the way of modeling

```
training2 <- training

for(x in ncol(training2):1) {
  if(anyNA(training2[,x]))
    training2[,x] <- NULL
}

training2ColNames <- colnames(training2)

for(x in ncol(training2):1) {
  #print(x)
  strTest <- substr(training2ColNames[x], 1,4)
  if(strTest == "kurt" | strTest == "skew" | strTest == "min_" | strTest == "max_" |
strTest == "ampl")
    training2[,x] <- NULL
}

training2$X<-NULL
training2$user_name<-NULL
training2$cvtd_timestamp<-NULL
training2$new_window<-NULL
```

Pre- processing


PCA analysis was performed and it was determined that while 2 PCA variables account for 99.9% of variance but there isn't a clear segregation of the data (see scatter Below) so I didn't continue with a PCA approach

```
PCA <- prcomp(training2[,-56])
summary(PCA)
```

```
## Importance of components:
```

```
##           PC1      PC2    PC3    PC4    PC5    PC6    PC7
## Standard deviation 2.882e+05 2.049e+05 599.4 526.3 440.6 358.6 311.3
## Proportion of Variance 6.642e-01 3.358e-01 0.0 0.0 0.0 0.0 0.0
## Cumulative Proportion 6.642e-01 1.000e+00 1.0 1.0 1.0 1.0 1.0
##           PC8    PC9   PC10   PC11   PC12   PC13   PC14   PC15
## Standard deviation 252.6 232.8 196.1 172.5 143.4 117.6 90.86 76.31
## Proportion of Variance 0.0 0.0 0.0 0.0 0.0 0.0 0.00 0.00
## Cumulative Proportion 1.0 1.0 1.0 1.0 1.0 1.0 1.00 1.00
##           PC16 PC17   PC18   PC19   PC20   PC21 PC22   PC23   PC24
## Standard deviation 70.13 62.5 58.02 53.48 52.28 49.58 42.9 39.47 36.39
## Proportion of Variance 0.00 0.0 0.00 0.00 0.00 0.00 0.0 0.00 0.00
## Cumulative Proportion 1.00 1.0 1.00 1.00 1.00 1.00 1.0 1.00 1.00
##           PC25   PC26   PC27   PC28   PC29   PC30   PC31   PC32
## Standard deviation 33.83 32.52 27.81 25.08 23.25 21.47 18.98 17.12
## Proportion of Variance 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00
## Cumulative Proportion 1.00 1.00 1.00 1.00 1.00 1.00 1.00 1.00
##           PC33   PC34   PC35   PC36   PC37   PC38   PC39   PC40
## Standard deviation 14.99 13.98 9.752 7.431 6.957 6.674 6.124 4.118
## Proportion of Variance 0.00 0.00 0.000 0.000 0.000 0.000 0.000 0.000
## Cumulative Proportion 1.00 1.00 1.000 1.000 1.000 1.000 1.000 1.000
##           PC41   PC42   PC43   PC44   PC45   PC46   PC47   PC48
## Standard deviation 3.644 3.471 3.314 1.95 1.487 1.078 0.4626 0.3933
## Proportion of Variance 0.000 0.000 0.000 0.00 0.000 0.000 0.0000 0.0000
## Cumulative Proportion 1.000 1.000 1.000 1.00 1.000 1.000 1.0000 1.0000
##           PC49   PC50   PC51   PC52   PC53   PC54   PC55
## Standard deviation 0.359 0.3143 0.2409 0.2019 0.1859 0.1038 0.0368
## Proportion of Variance 0.000 0.0000 0.0000 0.0000 0.0000 0.0000 0.0000
## Cumulative Proportion 1.000 1.0000 1.0000 1.0000 1.0000 1.0000 1.0000
```

```
typeColor <- ifelse(training$classe == "A",1, ifelse(training$classe == "B", 2, ifelse(
training$classe == "C", 3,ifelse(training$classe == "D",4,5))))
```



Model Selection

Prior to model selection I chose to center and scale the data. I first attempted a glm model but it failed to converge. Next I used a RF and was able to complete and predict 100% on the training set I partitioned as a cross reference check. Code and results are detailed below

Step 1 Fit RF

```
pproc <- preProcess(training2[, -56], method=c("center", "scale"))
training3 <- predict(pproc, training2)
modFit <- train(classe~., method="rf", data=training3)
```

Partition / subset training data with by creating a index at random to set aside for a cross reference approach

```
randomSampleIndex <- sample(1:nrow(training3))
training4 <- training3[randomSampleIndex,]
pred4 <- predict(modFit, training4)

crossRefComp <- data.frame(pred4, training4$classe)
numErrors <- ifelse(pred4 == training4$classe, 0, 1)
```

Number of Errors after predicting the cross reference training data is equal to 0

Model evaluation

Looking at the confusion matrix shows that we have a very small error rat of just .06%

```
modFit$finalModel

##
## Call:
##  randomForest(x = x, y = y, mtry = param$mtry)
##              Type of random forest: classification
##              Number of trees: 500
## No. of variables tried at each split: 28
##
##              OOB estimate of  error rate: 0.06%
## Confusion matrix:
##      A      B      C      D      E  class.error
## A 5580      0      0      0      0 0.0000000000
## B      2 3794      1      0      0 0.0007900974
## C      0      4 3418      0      0 0.0011689071
## D      0      0      3 3212      1 0.0012437811
## E      0      0      0      1 3606 0.0002772387
```

With a strong model we move pre-process and the test data to evaluate the model results

Prepare Test data and run a prediction

```

testing2 <- testing

for(x in ncol(testing2):1) {
  if(anyNA(testing2[,x]))
    testing2[,x] <- NULL
}

testing2ColNames <- colnames(testing2)

for(x in ncol(testing2):1) {

  strTest <- substr(testing2ColNames[x], 1,4)
  if(strTest == "kurt" | strTest == "skew" | strTest == "min_" | strTest == "max_" |
strTest == "ampl")
    testing2[,x] <- NULL
}

testing2$X<-NULL
testing2$user_name<-NULL
testing2$cvtd_timestamp<-NULL
testing2$new_window<-NULL

testing3 <- predict(pproc, testing2)
predTest <- predict(modFit, testing3)

```

Result / Conclusion

The prediction came in at 100% accurate based on seperate quiz answers

Predictions for the test set is as follows

```

##    [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E

```